

# Lecture 12: Neural Networks

Sergei V. Kalinin

# Types of Machine Learning

## **Supervised (inductive) learning**

- Given: training data + desired outputs (labels)

## **• Unsupervised learning**

- Given: training data (without desired outputs)

## **• Semi-supervised learning**

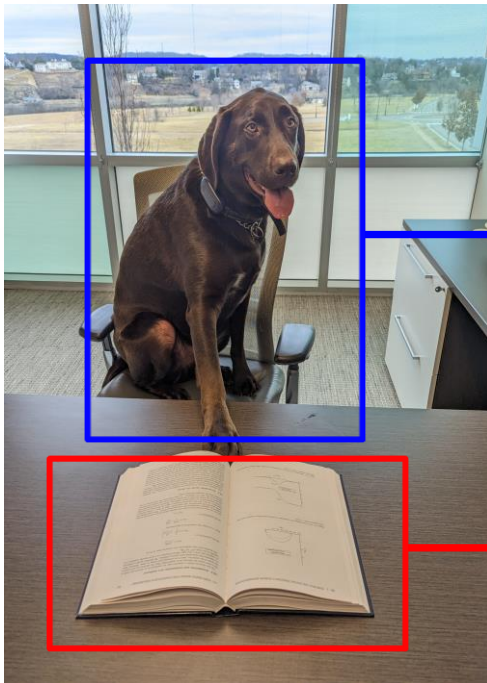
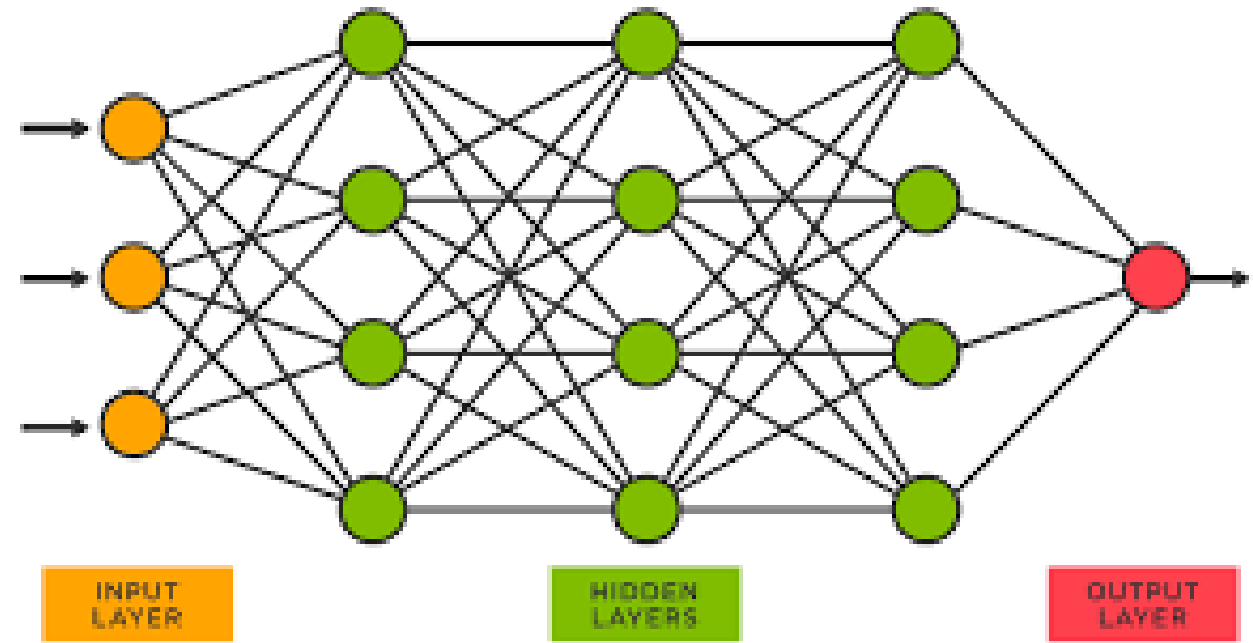
- Given: training data + a few desired outputs

## **• Reinforcement learning**

- Rewards from sequence of actions

# Supervised Machine Learning

- Regression
- Classification
- Semantic segmentation
- Instance segmentation
- ...



Dog

Book

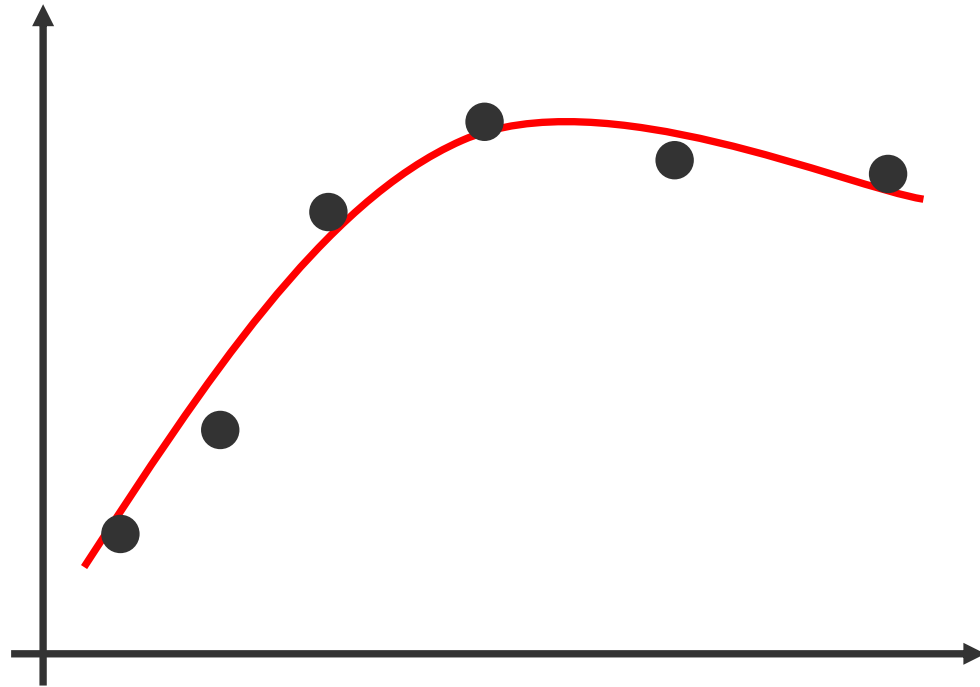
# Classification

- Given  $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- Learn a function  $f(x)$  to predict  $y$  given  $x$
- If  $y$  is categorical == classification

| Application                   | Input Data           | Classification                         |
|-------------------------------|----------------------|--|
| Medical Diagnosis             | Noninvasive tests    | Results from invasive measurements     |
| Optical Character Recognition | Scanned bitmaps      | Letter A-Z and digits 0-9              |
| Protein Folding               | Amino acid sequence  | Protein shape (helices, loops, sheets) |
| Materials Discovery           | Composition          | Metal/Semiconducotr                    |
| Research Paper Acceptance     | Words in paper title | Paper accepted or rejected             |

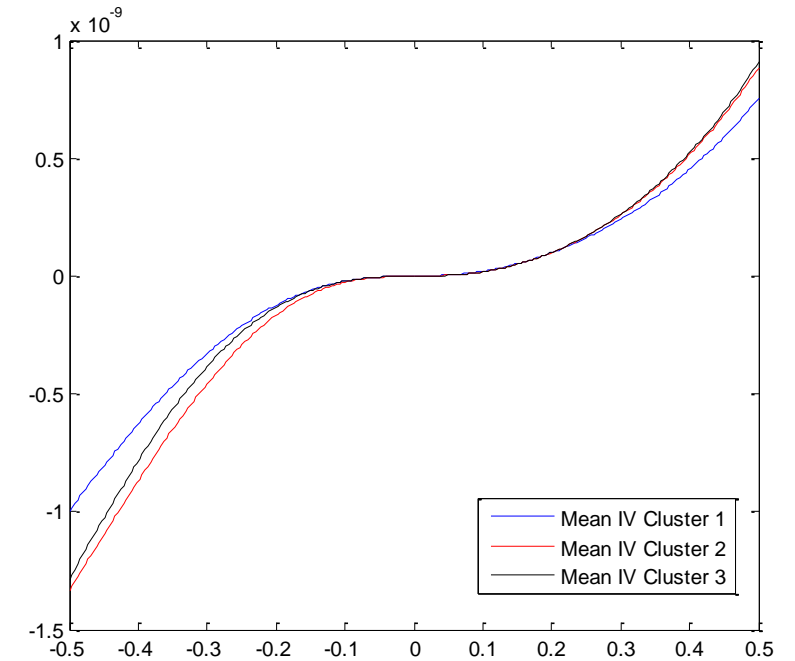
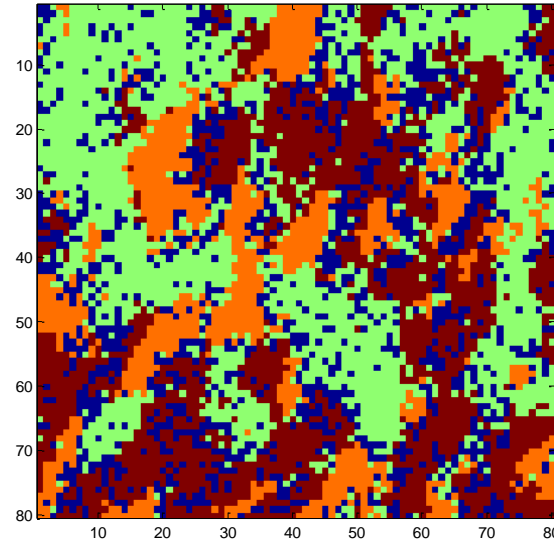
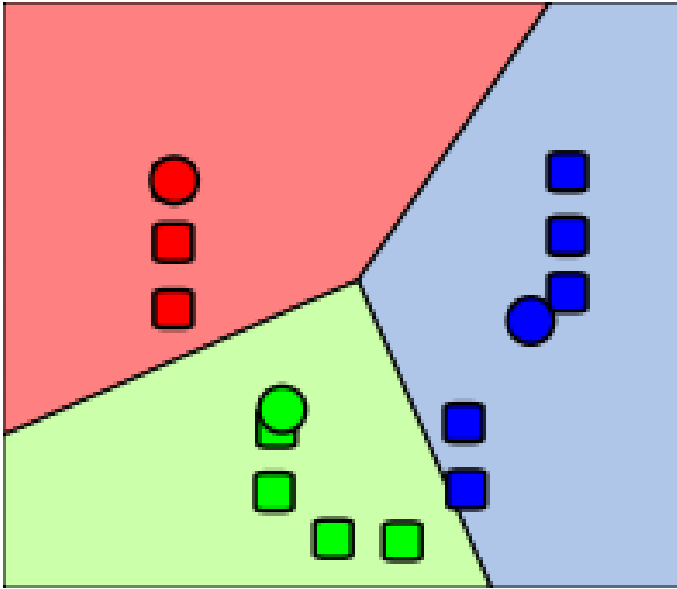
# Regression

- Given  $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- Learn a function  $f(x)$  to predict  $y$  given  $x$
- $y$  is real-valued == regression



# Unsupervised Learning

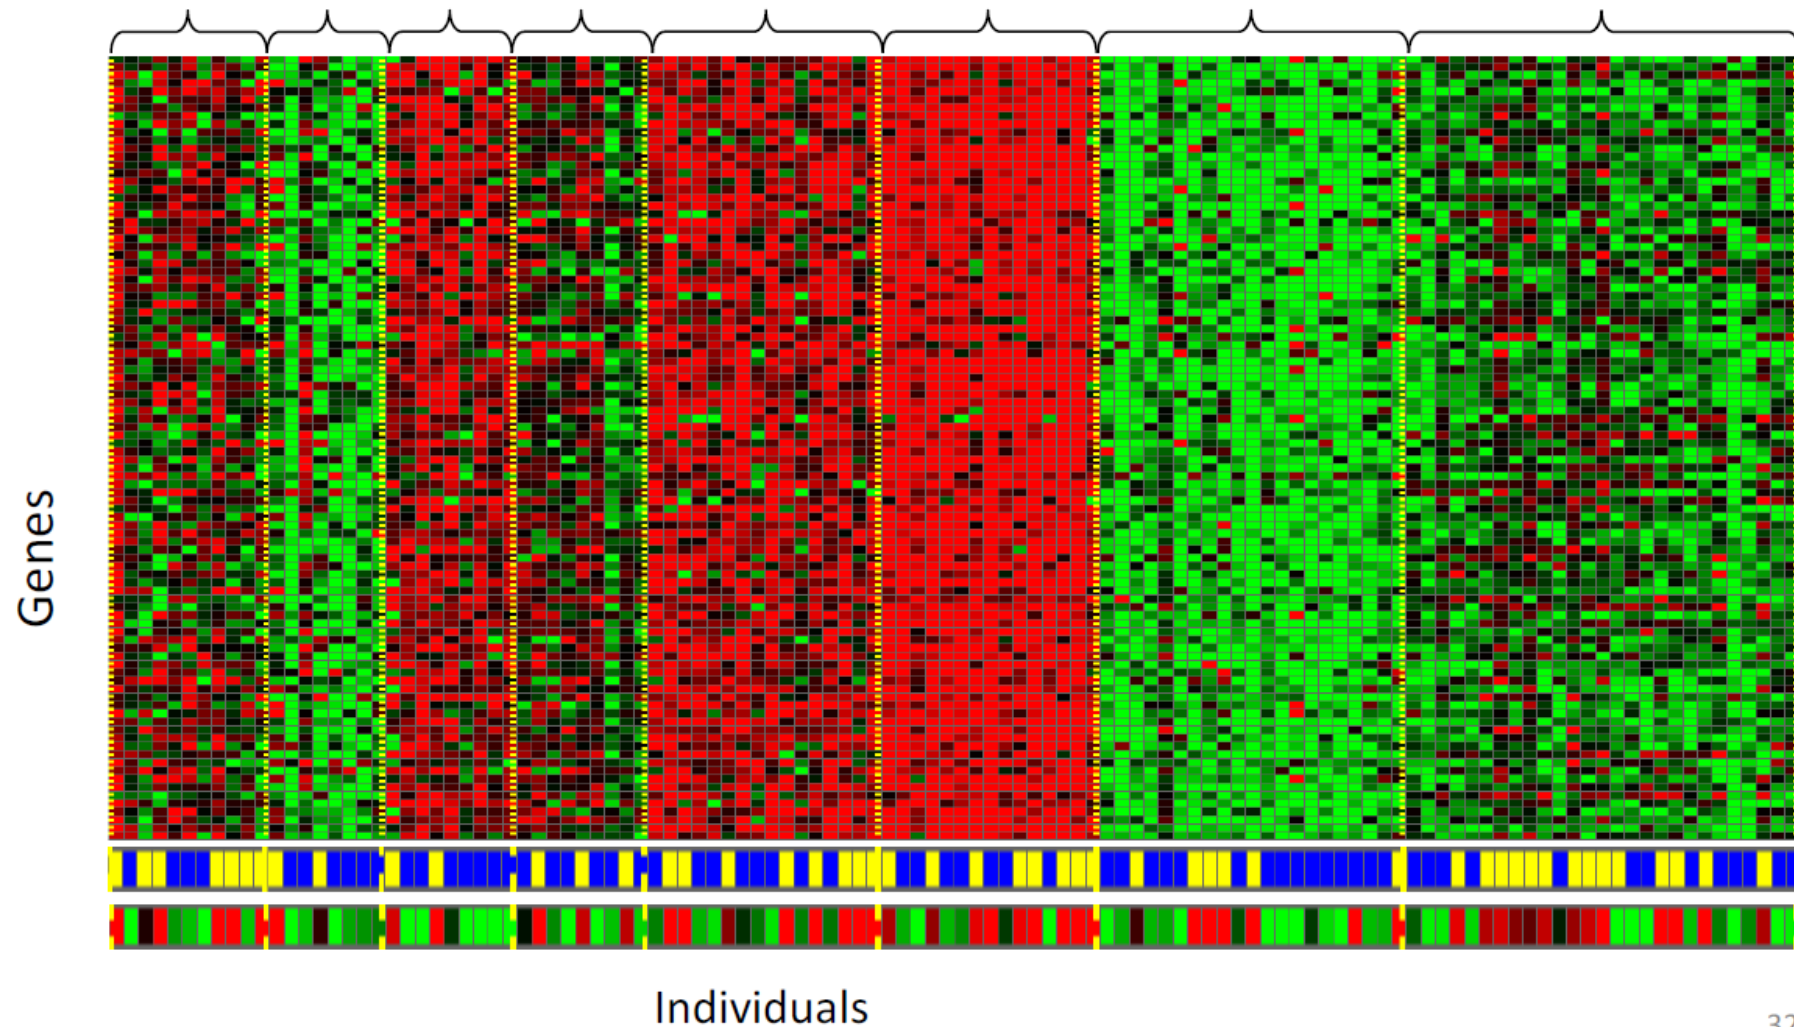
- Given  $x_1, x_2, \dots, x_n$  (without labels)
- Output hidden structure behind the  $x$ 's
- E.g., clustering



M. ZIATDINOV, A. MAKSOV, L. LI, A. SEFAT, P. MAKSYMОВYCH, and S.V. KALININ, *Deep data mining in a real space: Separation of intertwined electronic responses in a lightly-doped BaFe<sub>2</sub>As<sub>2</sub>*, Nanotechnology **27**, 475706 (2016).

# Unsupervised Learning

Genomics application: group individuals by genetic similarity

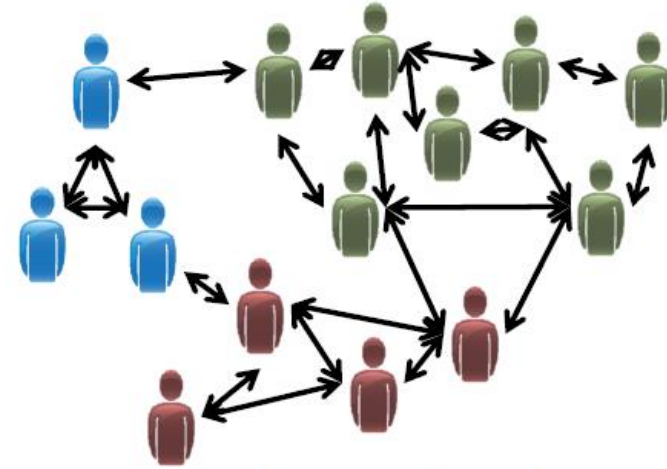


[Source: Daphne Koller]

# Unsupervised Learning



Organize computing clusters



Social network analysis



Market segmentation

Slide credit: Andrew Ng

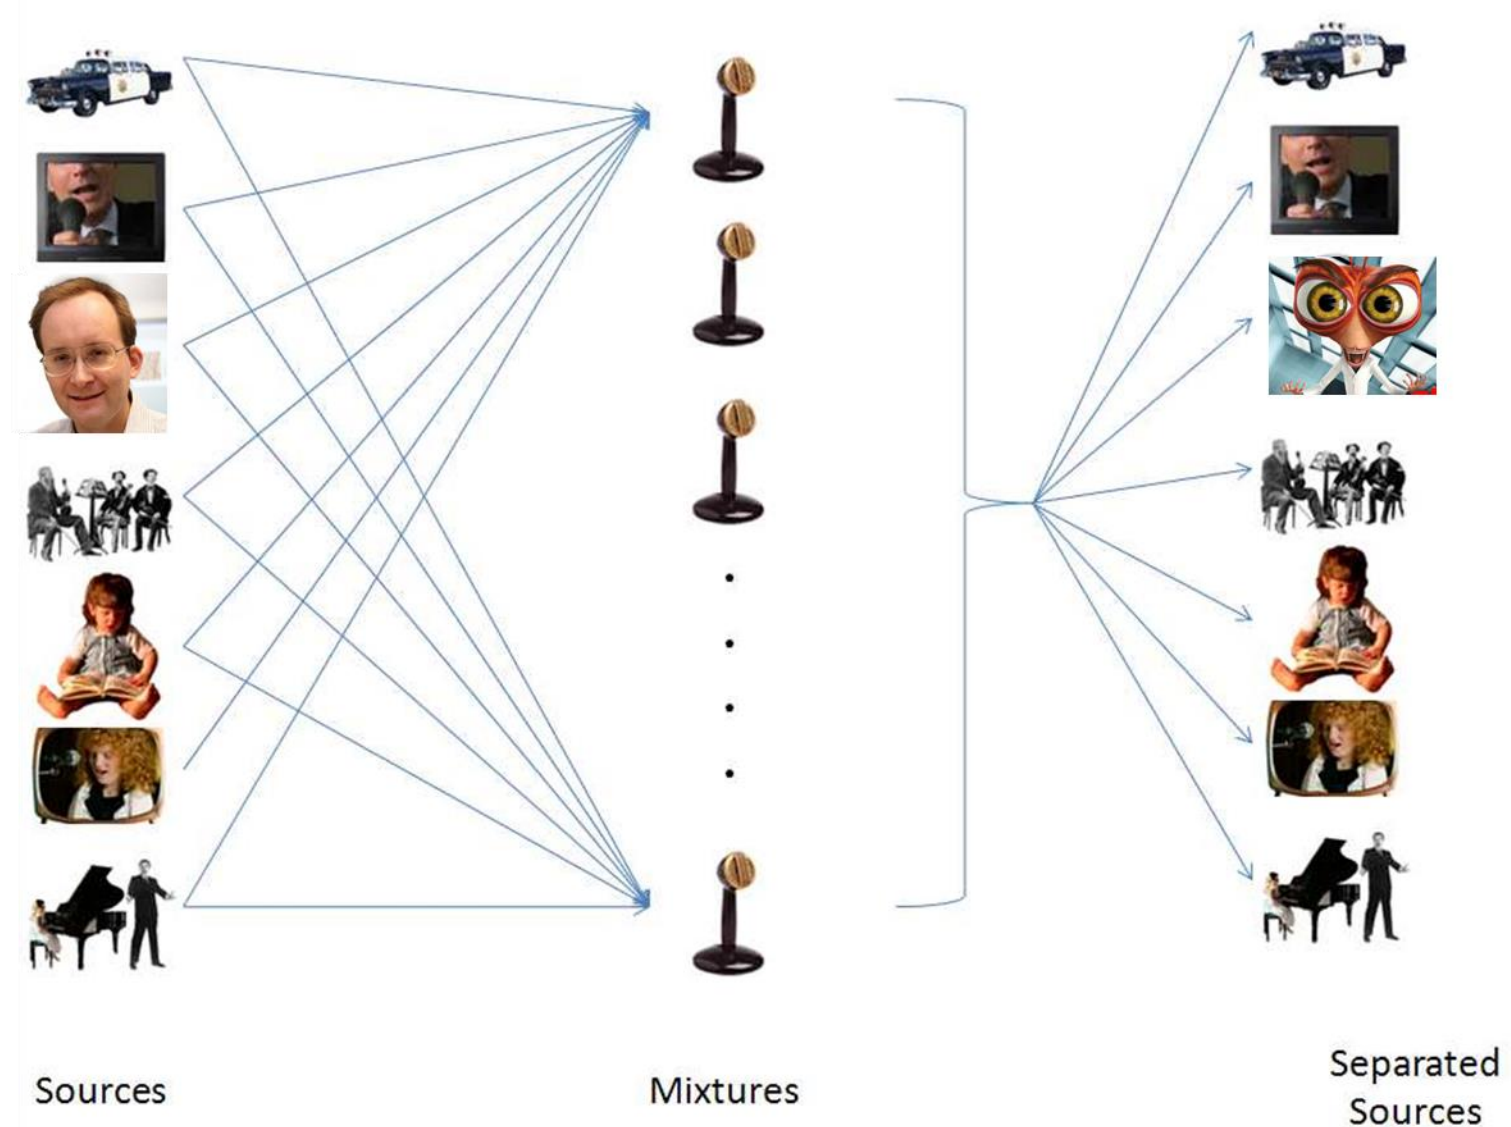


Astronomical data analysis



# Unsupervised Learning

Number of signals are being produced simultaneously; with the objective of separating and following each source separately



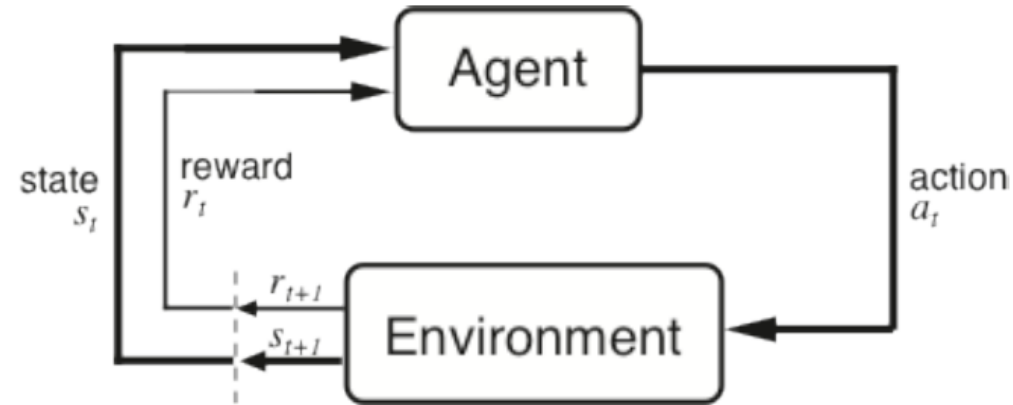
# Reinforcement Learning

Given a sequence of states and actions with (delayed) rewards, output a policy

- Policy is a mapping from states to actions that tells you what to do in a given state

- Examples:
  - Credit assignment problem
  - Game playing
  - Robot in a maze
  - Balance a pole on your hand

# RL: Agent and Environment



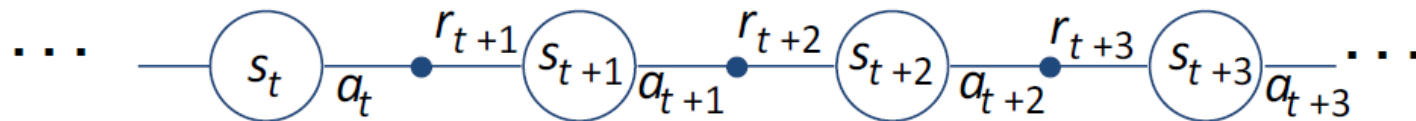
Agent and environment interact at discrete time steps :  $t = 0, 1, 2, K$

Agent observes state at step  $t$ :  $s_t \in S$

produces action at step  $t$ :  $a_t \in A(s_t)$

gets resulting reward :  $r_{t+1} \in \mathfrak{R}$

and resulting next state :  $s_{t+1}$



# Reinforcement Learning in Action

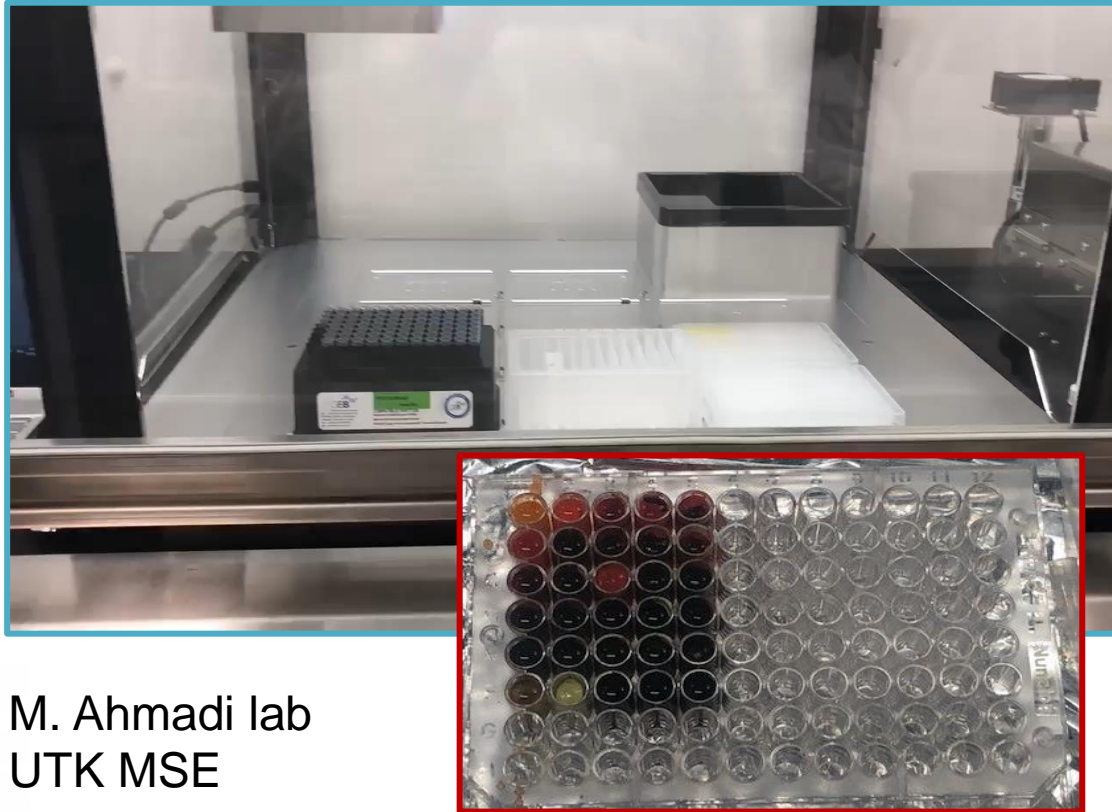


<https://www.youtube.com/watch?v=GtYIVxv0py8>



# Reinforcement Learning Applications

## Chemical Synthesis and Drug Discovery



M. Ahmadi lab  
UTK MSE

## Cloud Laboratories



Emerald Cloud Lab,  
SF and CMU

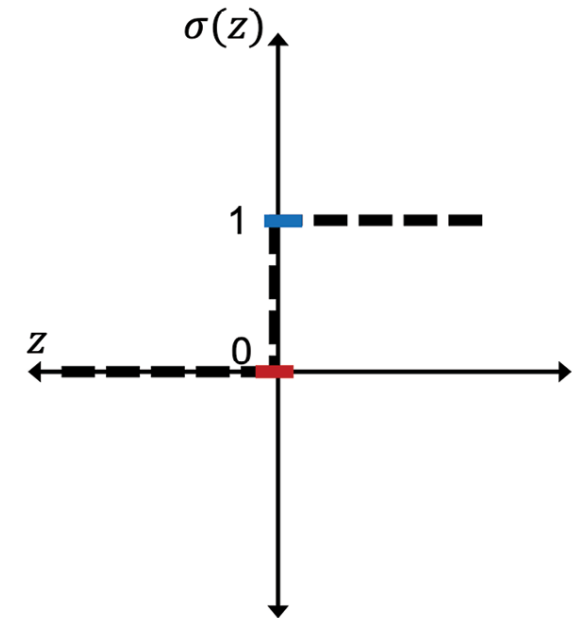
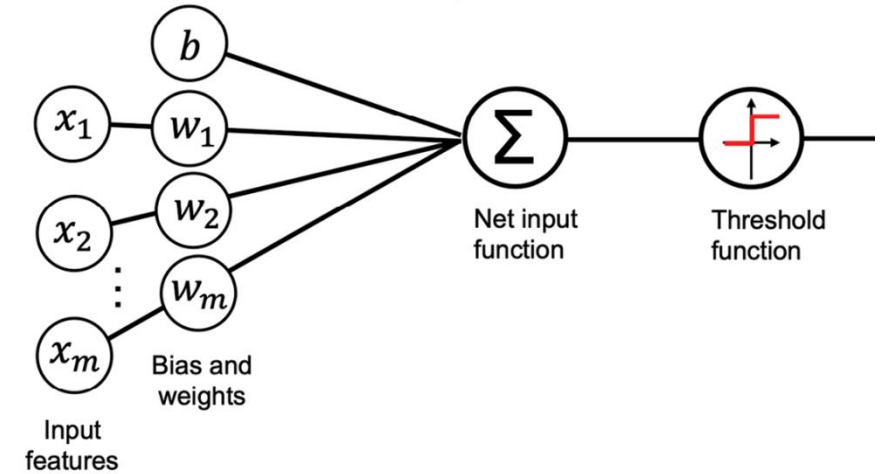
# Building Linear Neuron

Input:  $\mathbf{x} = \begin{bmatrix} x_1 \\ \vdots \\ x_m \end{bmatrix}$

Weights:  $\mathbf{w} = \begin{bmatrix} w_1 \\ \vdots \\ w_m \end{bmatrix}$

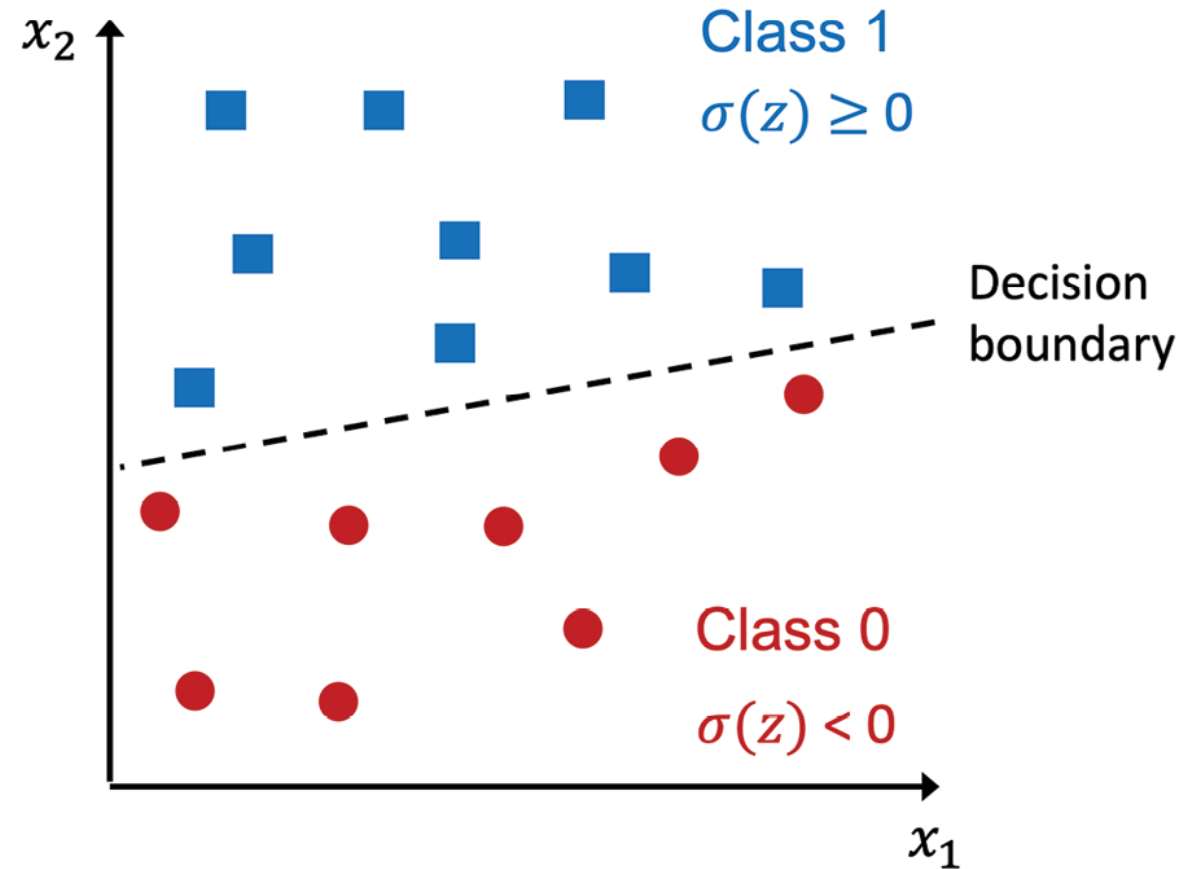
Linear transform: 
$$z = w_1x_1 + \dots + w_mx_m + b = \mathbf{w}^T\mathbf{x} + b$$

Output: 
$$\sigma(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{otherwise} \end{cases}$$



where  $z = \mathbf{w}^T\mathbf{x} + b$

# Linear Neuron in 2D



Linear transform:

$$z = w_1 x_1 + w_2 x_2 + b$$

Line:

$$x_2 = -w_1/w_2 x_1 - b/w_2$$

From S. Raschka, Machine Learning with PyTorch and Scikit-Learn

# Training Linear Neuron

- Initialize the weights and bias unit to 0 or small random numbers
- For each training example,  $\mathbf{x}(\mathbf{i})$ :
- Compute the output value,  $\mathbf{y}(\mathbf{i}) = \mathbf{w}^T \mathbf{x}(\mathbf{i}) + \mathbf{b}$
- Update the weights and bias unit:  $w_j := w_j + \Delta w_j$  and  $b := b + \Delta b$
- Where  $\Delta w_j = \eta(y^{(i)} - \hat{y}^{(i)})x_j^{(i)}$  and  $\Delta b = \eta(y^{(i)} - \hat{y}^{(i)})$

Each weight,  $w_j$ , corresponds to a feature,  $x_j$ , in the dataset,

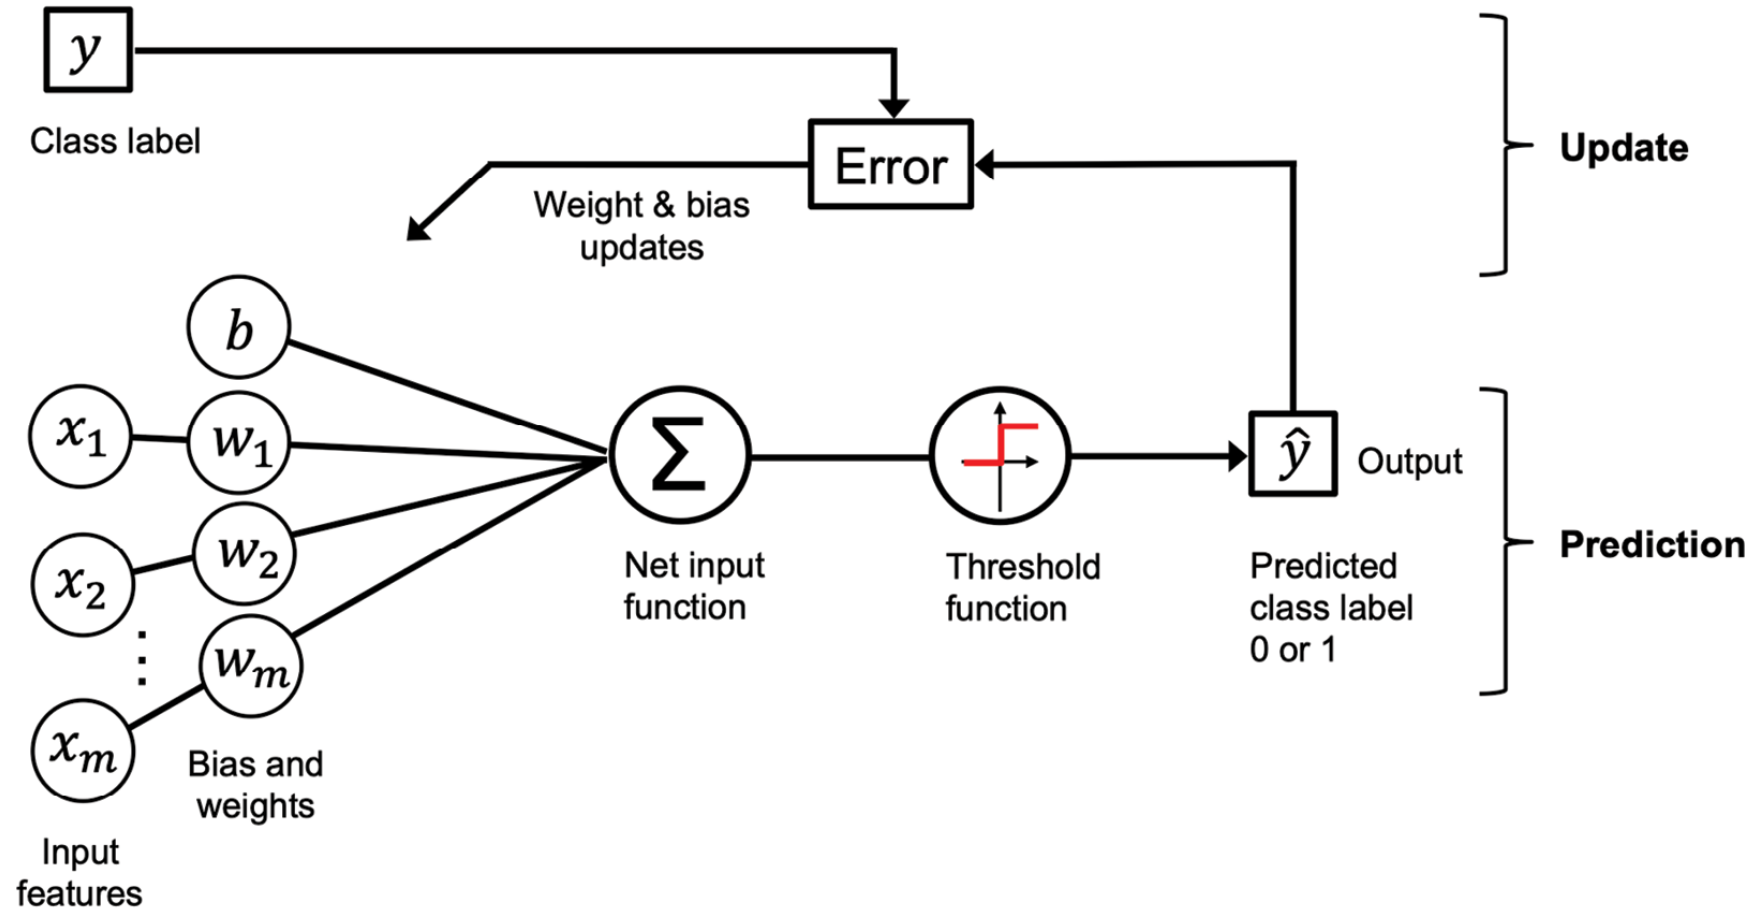
$\eta$  is the **learning rate** (typically a constant between 0.0 and 1.0),

$y^{(i)}$  is the **true class label** of the  $i$ th training example,

$\hat{y}^{(i)}$  is the **predicted class label**



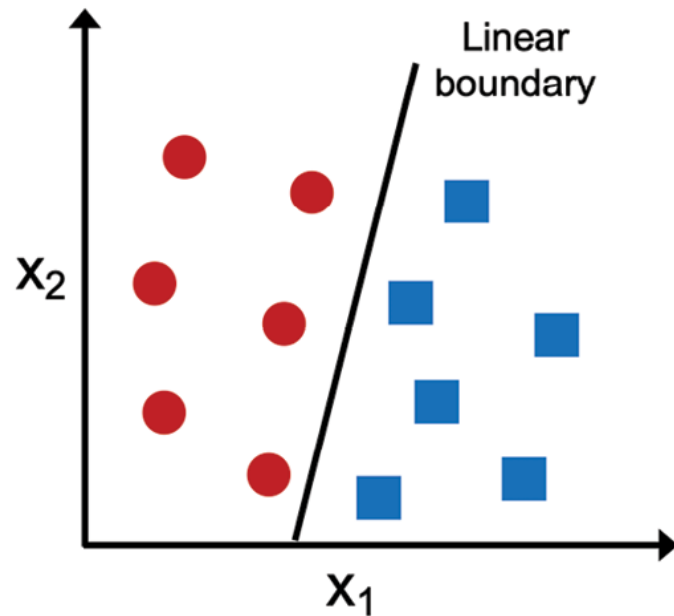
# Training Linear Neuron



# What problems can perceptron solve?

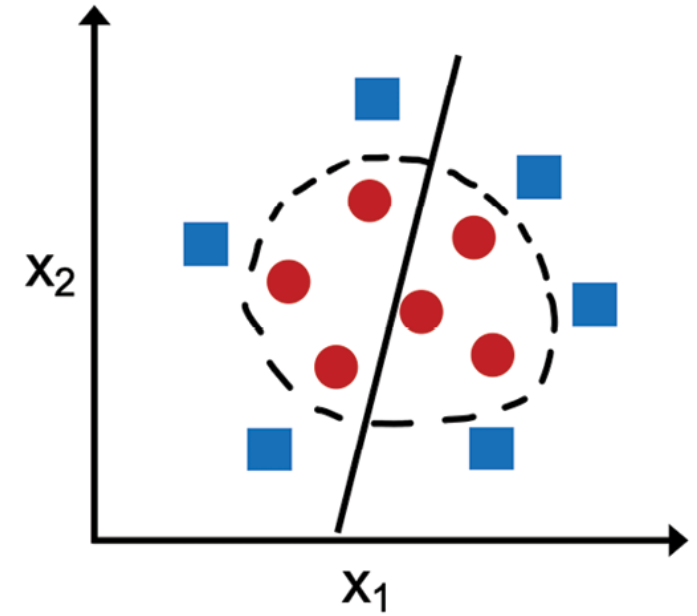
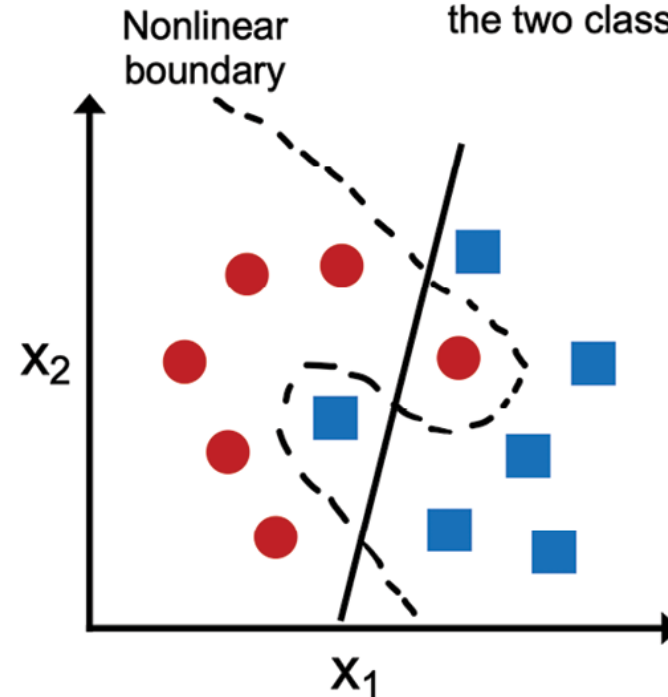
## Linearly separable

A linear decision boundary that separates the two classes exists

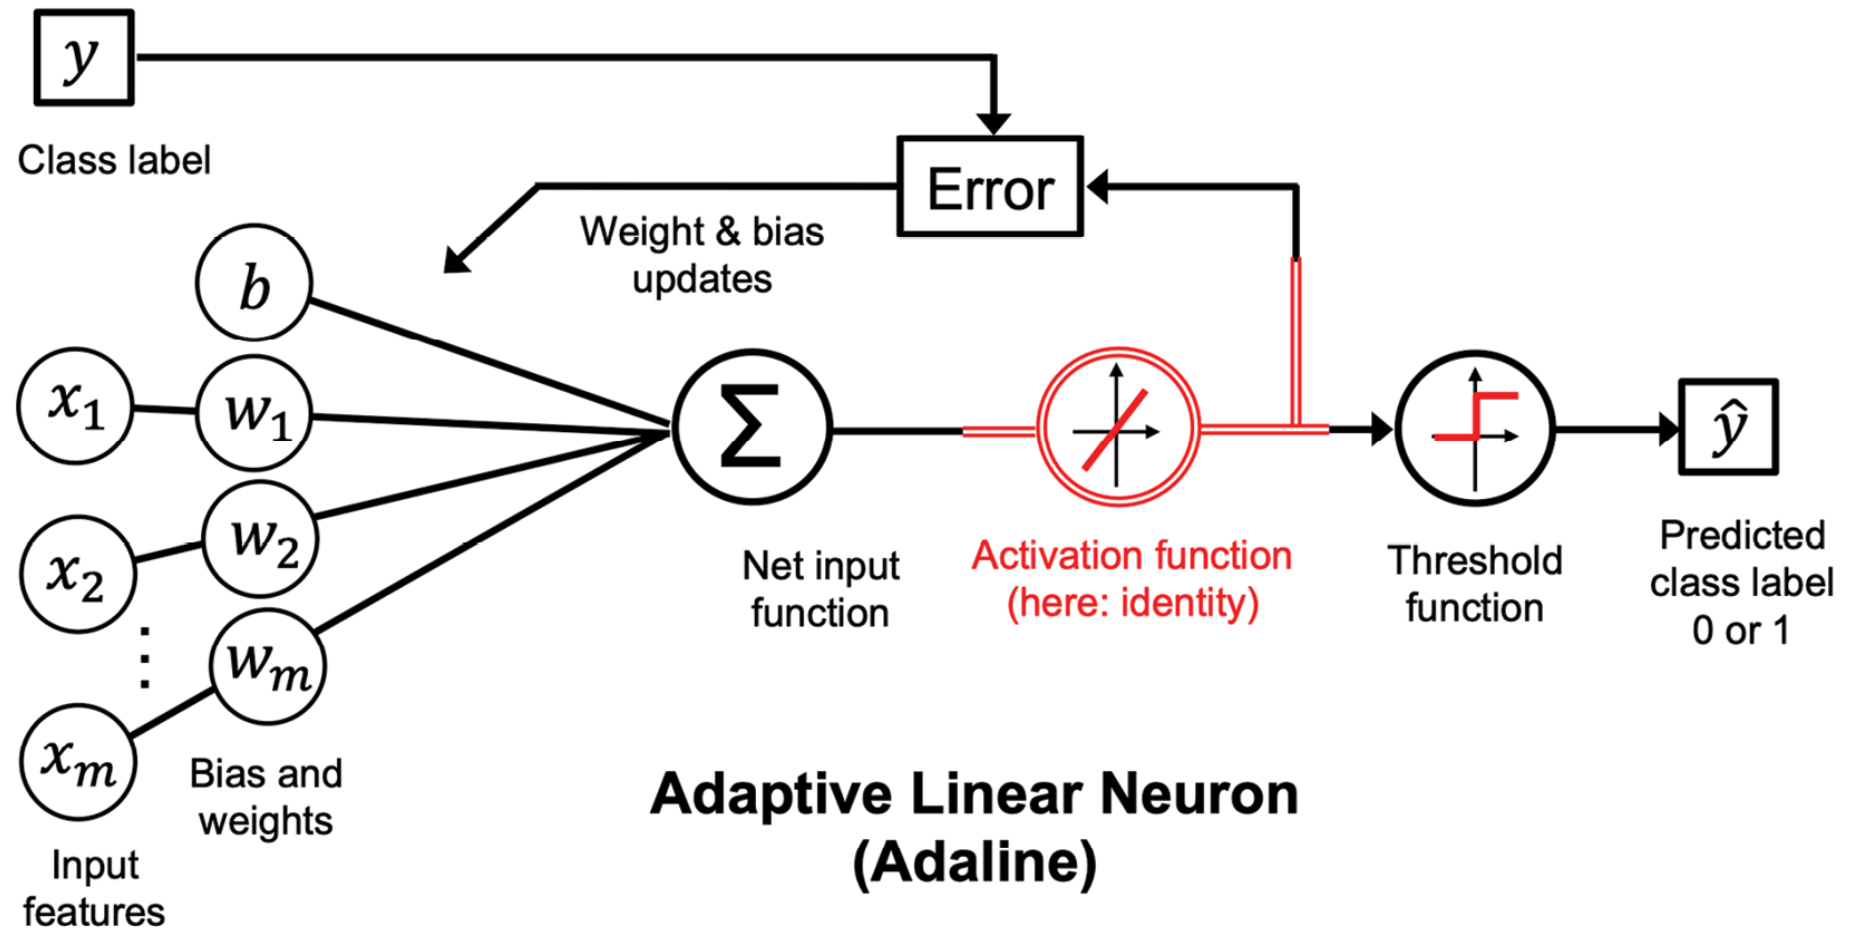


## Not linearly separable

No linear decision boundary that separates the two classes perfectly exists



# Adaline



# Adaline training

Loss function:

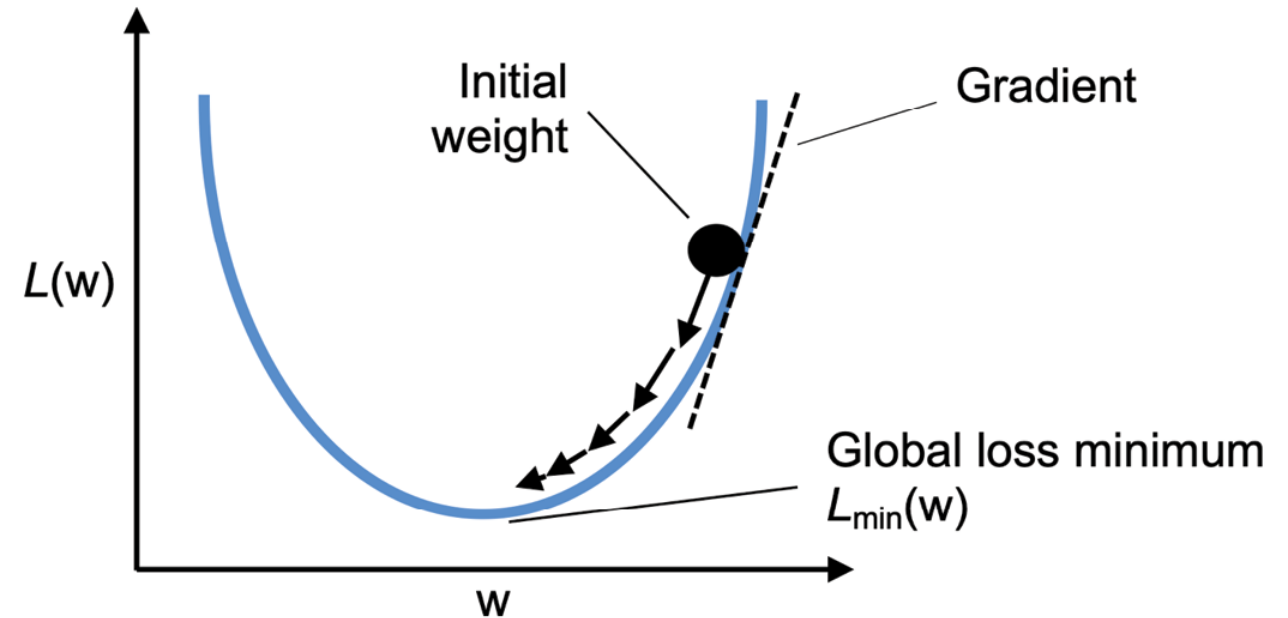
$$L(\mathbf{w}, b) = \frac{1}{n} \sum_{i=1}^n \left( y^{(i)} - \sigma(z^{(i)}) \right)^2$$

Weights update:

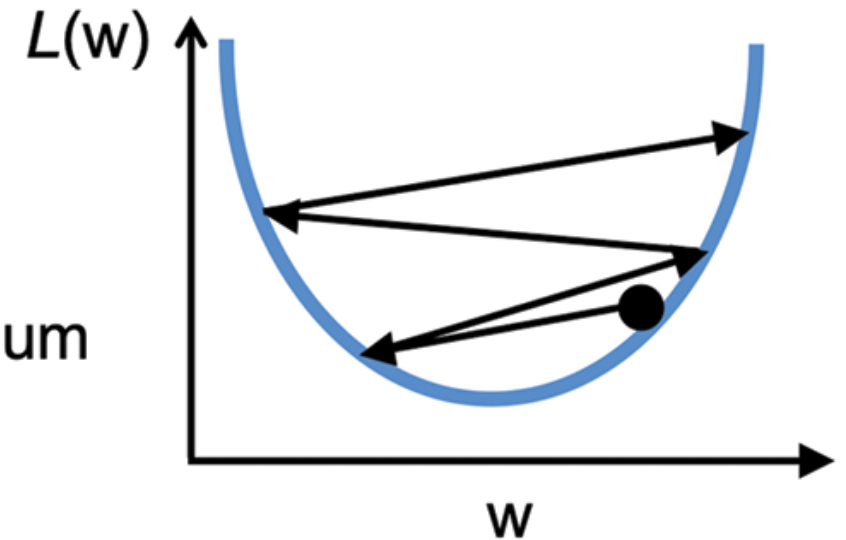
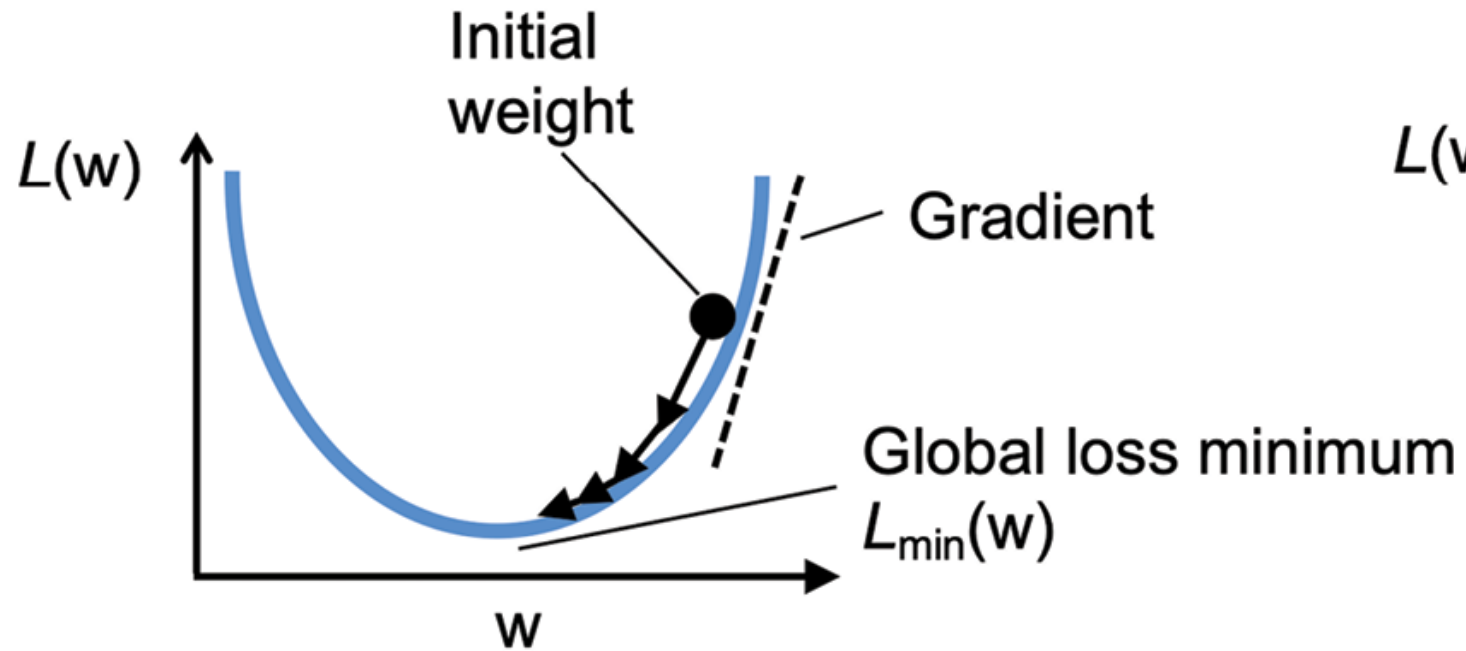
$$\mathbf{w} := \mathbf{w} + \Delta \mathbf{w}, \quad b := b + \Delta b$$

Learning rule:

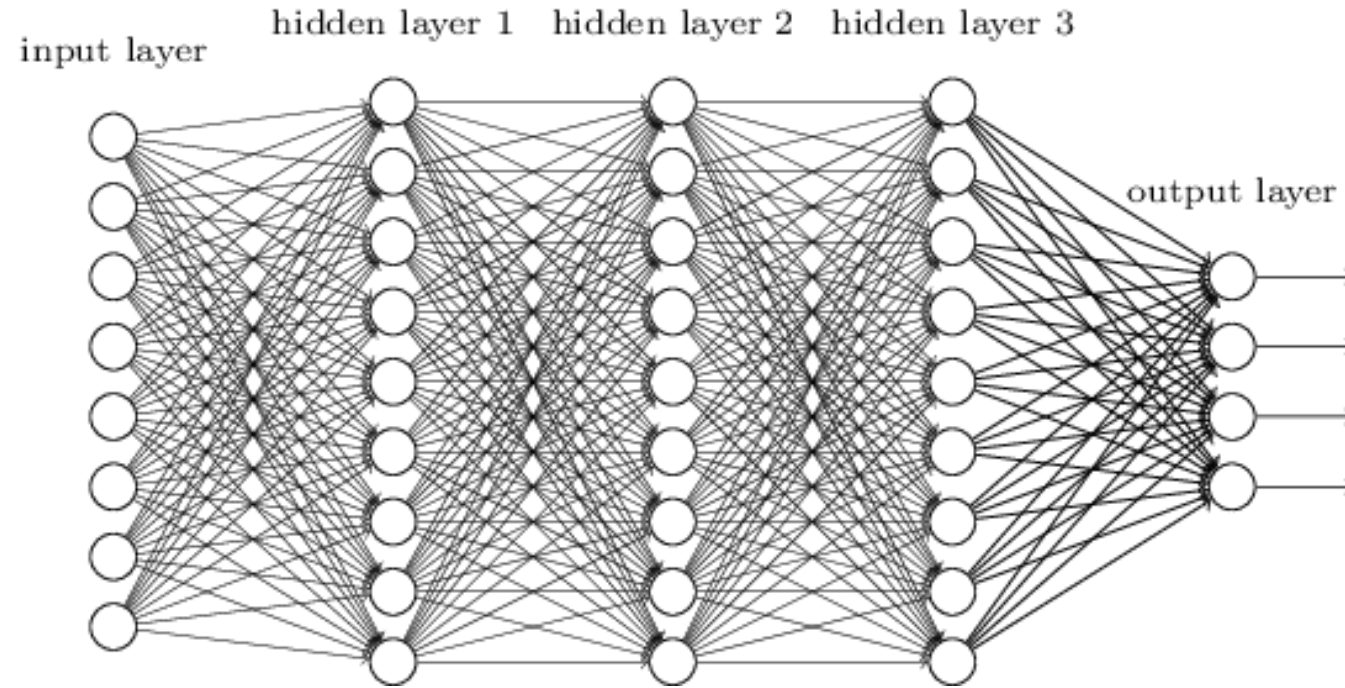
$$\Delta \mathbf{w} = -\eta \nabla_{\mathbf{w}} L(\mathbf{w}, b), \quad \Delta b = -\eta \nabla_b L(\mathbf{w}, b)$$



# Role of learning rate



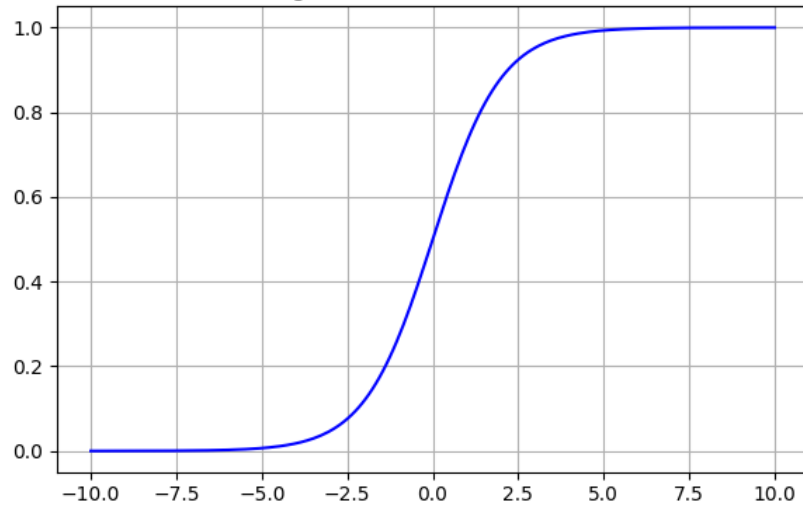
# Putting Neurons Together



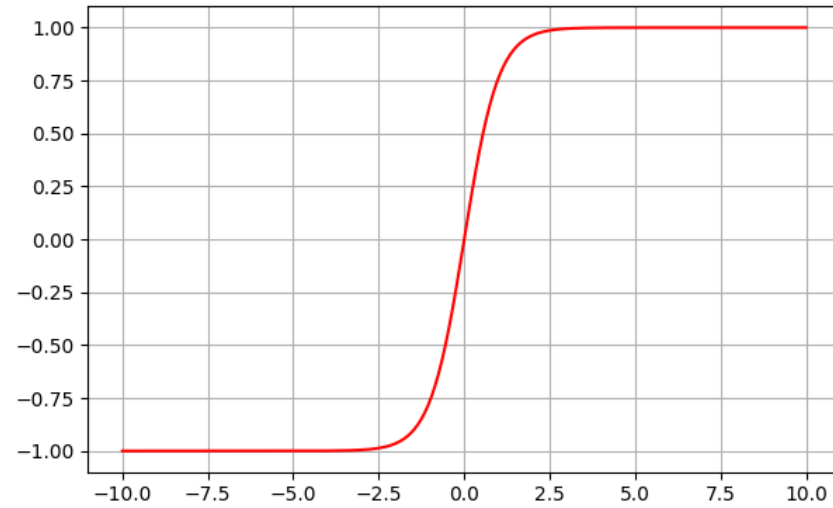
- Composed of multiple layers of artificial neurons.
- Each layer processes inputs received, applies a transformation (weights, biases, activation function), and passes the output to the next layer.
- Training a DNN involves adjusting weights and biases using backpropagation and a chosen optimization algorithm.
- The deep architecture enable the network to learn complex and abstract patterns in data.

# Activation functions

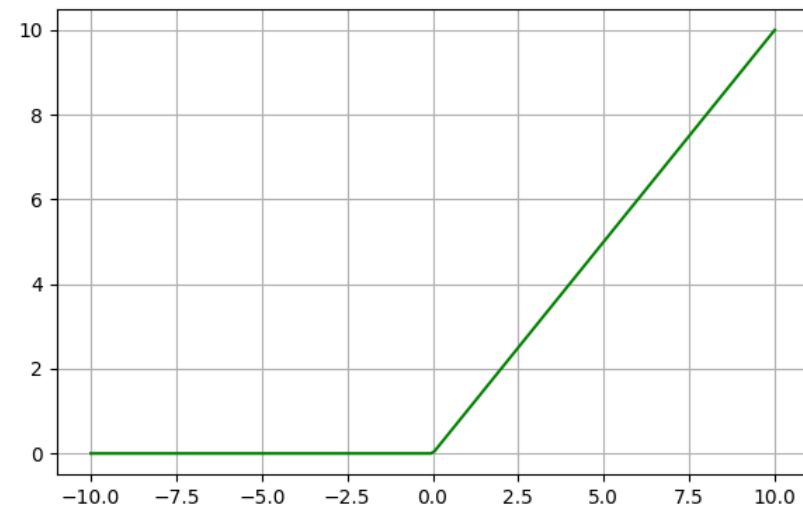
Sigmoid Activation Function



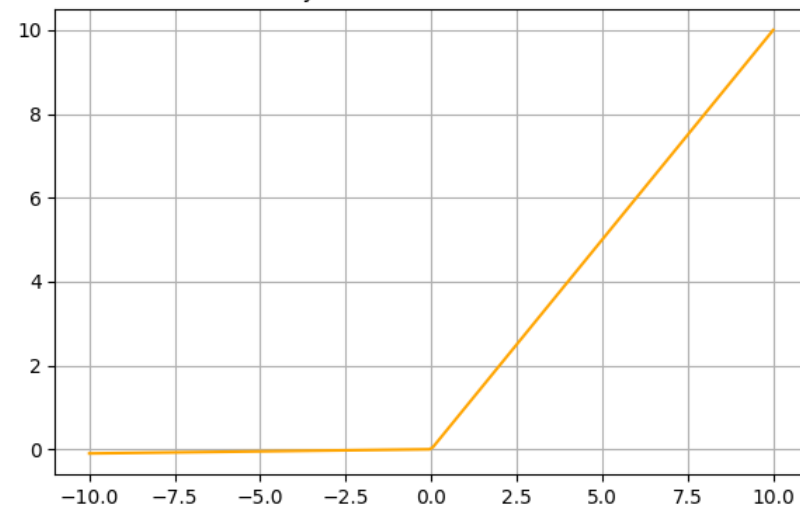
Tanh Activation Function



ReLU Activation Function



Leaky ReLU Activation Function



# Loss functions for supervised ML

A loss function, also known as a cost function, quantifies the difference between the predicted values and the actual target values. It guides the training of neural networks by providing a measure to minimize during optimization

**Mean Squared Error (MSE):** Used for regression problems.

$$MSE = \frac{1}{N} \sum_{i=1}^N (y_i - \hat{y}_i)^2$$

Measures the average squared difference between actual and predicted values

**Cross-Entropy Loss:** Used for classification problems.

$$CE = - \sum_{i=1}^N y_i \log(\hat{y}_i)$$

Measures the performance of a classification model whose output is a probability value between 0 and 1

- Loss functions provide the primary feedback signal for learning.
- The choice of loss function can significantly affect the model's performance and convergence



# Backpropagation

Backpropagation is a mechanism used to update the weights in a neural network efficiently, based on the error rate obtained in the previous epoch (i.e., iteration). It effectively distributes the error back through the network layers

- **Forward Pass:** Calculating the predicted output, moving the input data through the network layers
- **Loss Function:** Determining the error by comparing the predicted output to the actual output
- **Backward Pass:** Computing the gradient of the loss function with respect to each weight by the chain rule
- **Weight Update:** Adjusting the weights of the network in a direction that minimally reduces the loss (gradient descent)

**Input Data → Forward Pass → Calculate Loss → Backward Pass → Update Weights**

<https://medium.com/@14prakash/back-propagation-is-very-simple-who-made-it-complicated-97b794c97e5c>

# How to not get lost?

- A vast array of network architectures ranging from Multilayer Perceptrons (MLPs) to Graph Neural Networks and Transformers
- Each architecture has unique characteristics suited for different types of data and tasks, ways to define the architecture, and so on
- Numerous methods to engineer loss functions
- Numerous ways to implement regularization

## **Problem-Centric Approach:**

- Always start with the problem at hand: Analyze the nature of input data and desired output
- Choose a network architecture that aligns with the type and structure of your data
- Select a loss function that reflects the objective of the problem
- Metrics should be chosen based on what measures success for your specific task

## **Hyperparameter Tuning:**

- Once the architecture and loss function are set, proceed to tune hyperparameters including network structure, optimizers, regularization, etc.
- Hyperparameter tuning should be guided by the chosen metrics and the nature of the problem

## Some useful resources:

<https://udlbook.github.io/udlbook/>

<https://dmol.pub/ml/classification.html>

<https://keras.io/examples/>